

LALA

Autonomous desktop voice assistant · v1.0 -> v2.1 · MIT



Say "Hey Laala" - your desktop just listens. One brain, nine layers, zero paid dependencies.

PIPELINE

Wake -> Streaming STT -> Intent -> CEO/Planner -> Tools -> Response -> Streaming TTS (+barge-in)

THE NINE LAYERS

Multi-Agent Brain

CEO delegates to Research/Coding/Vision/Finance/Scheduler/Health/Home

World Model

network/battery/weather/calendar/processes/devices

Computer Vision

screen capture + OCR; webcam/face consent-ready

Voice Pipeline

streaming end-to-end, sub-second simple commands

Robotics

ESP32/Pi/arms/drones via HTTP; motion confirmed

Long-Term Memory

5 layers + consent-gated voiceprints

Autonomous Planning

interview prep: research->resume->questions->reminders

Digital Twin

live env mirror + what-changed diffs

Tool Ecosystem

terminal/code/email/calendar/maps/finance/home/PDF, gated

INSTALL - ONE CLICK

```
git clone https://github.com/1997mujtabaamjad-del/LALA.git && bash install.sh
```

Win: `iwr ../install.ps1 -UseBasicParsing | iex` Mac: `curl ../install-mac.sh | bash`

Single file: LALA-ALL-IN-ONE.html (app + effects + docs)

SELF-VERIFIED

54/54 functional checks 98 python tests 19 js tests secrets masked exec confirm-gated

bash scripts/autorun.sh · lala --milestone · lala --bench

LALA · intellect x enterprise · github.com/1997mujtabaamjad-del/LALA